

# CHAPTER 4



## Intermediate SQL

Solutions for the Practice Exercises of Chapter 4

### Practice Exercises

4.1

**Answer:**

Although the query is syntactically correct, it does not compute the expected answer because *dept\_name* is an attribute of both *course* and *instructor*. As a result of the natural join, results are shown only when an instructor teaches a course in her or his own department.

4.2

**Answer:**

- a. Display a list of all instructors, showing each instructor's ID and the number of sections taught. Make sure to show the number of sections as 0 for instructors who have not taught any section. Your query should use an outer join, and should not use subqueries.

```
select ID, count(sec_id) as Number_of_sections  
from instructor natural left outer join teaches  
group by ID
```

The above query should not be written using `count(*)` since that would count null values also. It could be written using any attribute from *teaches* which does not occur in *instructor*, which would be correct although it may be confusing to the reader. (Attributes that occur in *instructor* would not be null even if the instructor has not taught any section.)

- b. Write the same query as above, but using a scalar subquery, and not using outerjoin.

```

select ID,
       (select count(*) as Number_of_sections
        from teaches T where T.id = I.id)
from instructor I

```

- c. Display the list of all course sections offered in Spring 2018, along with the ID and name of each instructor teaching the section. If a section has more than one instructor, that section should appear as many times in the result as it has instructors. If a section does not have any instructor, it should still appear in the result with the instructor name set to “—”.

```

select course_id, sec_id, ID,
       decode(name, null, '—', name) as name
from (section natural left outer join teaches)
     natural left outer join instructor
where semester='Spring' and year= 2018

```

The query may also be written using the **coalesce** operator, by replacing **decode(..)** with **coalesce(name, '—')**. A more complex version of the query can be written using union of join result with another query that uses a subquery to find courses that do not match; refer to Exercise 4.3.

- d. Display the list of all departments, with the total number of instructors in each department, without using subqueries. Make sure to show departments that have no instructors, and list those departments with an instructor count of zero.

```

select dept_name, count(ID)
from department natural left outer join instructor
group by dept_name

```

### 4.3

#### Answer:

- a. **select \* from student natural left outer join takes**  
can be rewritten as:

```

select * from student natural join takes
union
select ID, name, dept_name, tot_cred, null, null, null, null, null
from student S1 where not exists
(select ID from takes T1 where T1.id = S1.id)

```

- b. **select \* from student natural full outer join takes**  
can be rewritten as:

```

(select * from student natural join takes)
union
(select ID, name, dept_name, tot_cred, null, null, null, null, null
from student S1
where not exists
  (select ID from takes T1 where T1.id = S1.id))
union
(select ID, null, null, null, course_id, sec_id, semester, year, grade
from takes T1
where not exists
  (select ID from student S1 where T1.id = S1.id))

```

## 4.4

## Answer:

- Consider  $r = (a, b)$ ,  $s = (b1, c1)$ ,  $t = (b, d)$ . The second expression would give  $(a, b, null, d)$ .
- Since  $s$  **natural left outer join**  $t$  is computed first, the absence of nulls in both  $s$  and  $t$  implies that each tuple of the result can have  $D$  null, but  $C$  can never be null.

## 4.5

## Answer:

- Consider the case where a professor in the Physics department teaches an Elec. Eng. course. Even though there is a valid corresponding entry in *teaches*, it is lost in the natural join of *instructor*, *teaches* and *course*, since the instructor's department name does not match the department name of the course. A dataset corresponding to the same is:

```

instructor = {'12345', 'Gauss', 'Physics', 10000}
teaches = {'12345', 'EE321', 1, 'Spring', 2017}
course = {'EE321', 'Magnetism', 'Elec. Eng.', 6}

```

- The query in question 4.2(a) is a good example for this. Instructors who have not taught a single course should have number of sections as 0 in the query result. (Many other similar examples are possible.)
- Consider the query

```
select * from teaches natural join instructor;
```

In this query, we would lose some sections if *teaches.ID* is allowed to be *null* and such tuples exist. If, just because *teaches.ID* is a foreign key to *instructor*, we did not create such a tuple, the error in the above query would not be detected.

## 4.6

**Answer:**

We should not add credits for courses with a null grade; further, to correctly handle the case where a student has not completed any course, we should make sure we don't divide by zero, and should instead return a null value.

We break the query into a subquery that finds sum of credits and sum of credit-grade-points, taking null grades into account. The outer query divides the above to get the average, taking care of divide by zero.

```
create view student_grades(ID, GPA) as
  select ID, credit_points / decode(credit_sum, 0, null, credit_sum)
from ((select ID, sum(decode(grade, null, 0, credits)) as credit_sum,
        sum(decode(grade, null, 0, credits*points)) as credit_points
        from(takes natural join course) natural left outer join grade_points
        group by ID)
union
select ID, null, null
from student
where ID not in (select ID from takes)
```

The view defined above takes care of *null* grades by considering the credit points to be 0 and not adding the corresponding credits in *credit\_sum*.

The query above ensures that a student who has not taken any course with non-null credits, and has *credit\_sum* = 0 gets a GPA of *null*. This avoids the division by zero, which would otherwise have resulted.

In systems that do not support **decode**, an alternative is the **case** construct. Using **case**, the solution would be written as follows:

```
create view student_grades(ID, GPA) as
  select ID, credit_points / (case when credit_sum = 0 then null
                             else credit_sum end)
from ((select ID, sum (case when grade is null then 0
                             else credits end) as credit_sum,
        sum (case when grade is null then 0
               else credits*points end) as credit_points
        from(takes natural join course) natural left outer join grade_points
        group by ID)
union
select ID, null, null
from student
where ID not in (select ID from takes)
```

---

```
create table employee
(ID          numeric(6,0),
 person_name char(20),
 street      char(30),
 city        char(30),
 primary key (ID))
```

```
create table works
(ID          numeric(6,0),
 company_name char(15),
 salary      integer,
 primary key (ID),
 foreign key (ID) references employee,
 foreign key (company_name) references company)
```

```
create table company
(company_name char(15),
 city        char(30),
 primary key (company_name))
```

```
create table manages
(ID          numeric(6,0),
 manager_iid numeric(6,0),
 primary key (ID),
 foreign key (ID) references employee,
 foreign key (manager_iid) references employee(ID))
```

---

**Figure 4.101** Figure for Exercise 4.7.

An alternative way of writing the above query would be to use *student* **natural left outer join** *gpa*, in order to consider students who have not taken any course.

#### 4.7

**Answer:**

Please see Figure 4.101.

Note that alternative data types are possible. Other choices for **not null** attributes may be acceptable.

## 4.8

**Answer:**

a. Query:

```

select  ID, name, sec_id, semester, year, time_slot_id,
        count(distinct building, room_number)
from    instructor natural join teaches natural join section
group by (ID, name, sec_id, semester, year, time_slot_id)
having  count(building, room_number) > 1

```

Note that the **distinct** keyword is required above. This is to allow two different sections to run concurrently in the same time slot and are taught by the same instructor without being reported as a constraint violation.

b. Query:

```

create assertion check not exists
( select ID, name, sec_id, semester, year, time_slot_id,
        count(distinct building, room_number)
from    instructor natural join teaches natural join section
group by (ID, name, sec_id, semester, year, time_slot_id)
having  count(building, room_number) > 1)

```

## 4.9

**Answer:**

The tuples of all employees of the manager, at all levels, get deleted as well! This happens in a series of steps. The initial deletion will trigger deletion of all the tuples corresponding to direct employees of the manager. These deletions will in turn cause deletions of second-level employee tuples, and so on, till all direct and indirect employee tuples are deleted.

## 4.10

**Answer:**

```

select coalesce(a.name, b.name) as name,
       coalesce(a.address, b.address) as address,
       a.title,
       b.salary
from   a full outer join b on a.name = b.name and
       a.address = b.address

```

## 4.11

**Answer:**

There are many reasons—we list a few here. One might wish to allow a user only to append new information without altering old information. One might wish

to allow a user to access a relation but not change its schema. One might wish to limit access to aspects of the database that are not technically data access but instead impact resource utilization, such as creating an index.

4.12

**Answer:**

Both cases give the same authorization at the time the statement is executed, but the long-term effects differ. If the grant is done based on the role, then the grant remains in effect even if the user who performed the grant leaves and that user's account is terminated. Whether that is a good or bad idea depends on the specific situation, but usually granting through a role is more consistent with a well-run enterprise.

4.13

**Answer:**

- No. This allows a user to be granted access to only part of relation  $r$ .
- Yes. A valid update issued using view  $v$  must update  $r$  for the update to be stored in the database.
- Any tuple  $t$  compatible with the schema for  $v$  but not satisfying the **where** clause in the definition of view  $v$  is a valid example. One such example appears in Section 4.2.4.

