

# CHAPTER 16



## Query Optimization

Solutions for the Practice Exercises of Chapter 16

### Practice Exercises

16.1

**Answer:**

The answer depends on the database.  
FILL IN

16.2

**Answer:**

a.  $E_1 \bowtie_0 (E_2 - E_3) = (E_1 \bowtie_0 E_2 - E_1 \bowtie_0 E_3).$

Let us rename  $(E_1 \bowtie_0 (E_2 - E_3))$  as  $R_1$ ,  $(E_1 \bowtie_0 E_2)$  as  $R_2$  and  $(E_1 \bowtie_0 E_3)$  as  $R_3$ . It is clear that if a tuple  $t$  belongs to  $R_1$ , it will also belong to  $R_2$ . If a tuple  $t$  belongs to  $R_3$ ,  $t[E_3\text{'s attributes}]$  will belong to  $E_3$ , hence  $t$  cannot belong to  $R_1$ . From these two we can say that

$$\forall t, t \in R_1 \Rightarrow t \in (R_2 - R_3)$$

It is clear that if a tuple  $t$  belongs to  $R_2 - R_3$ , then  $t[R_2\text{'s attributes}] \in E_2$  and  $t[R_2\text{'s attributes}] \notin E_3$ . Therefore:

$$\forall t, t \in (R_2 - R_3) \Rightarrow t \in R_1$$

The above two equations imply the given equivalence.

This equivalence is helpful because evaluation of the right-hand side join will produce many tuples which will finally be removed from the result. The left-hand side expression can be evaluated more efficiently.

- b.  $\sigma_\theta({}_A\gamma_F(E)) = {}_A\gamma_F(\sigma_\theta(E))$ , where  $\theta$  uses only attributes from  $A$ .

$\theta$  uses only attributes from  $A$ . Therefore if any tuple  $t$  in the output of  ${}_A\gamma_F(E)$  is filtered out by the selection of the left-hand side, all the tuples in  $E$  whose value in  $A$  is equal to  $t[A]$  are filtered out by the selection of the right-hand side. Therefore:

$$\forall t, t \notin \sigma_\theta({}_A\gamma_F(E)) \Rightarrow t \notin {}_A\gamma_F(\sigma_\theta(E))$$

Using similar reasoning, we can also conclude that

$$\forall t, t \notin {}_A\gamma_F(\sigma_\theta(E)) \Rightarrow t \notin \sigma_\theta({}_A\gamma_F(E))$$

The above two equations imply the given equivalence.

This equivalence is helpful because evaluation of the right-hand side avoids performing the aggregation on groups which are going to be removed from the result. Thus the right-hand side expression can be evaluated more efficiently than the left-hand side expression.

- c.  $\sigma_\theta(E_1 \bowtie E_2) = \sigma_\theta(E_1) \bowtie E_2$  where  $\theta$  uses only attributes from  $E_1$ .

$\theta$  uses only attributes from  $E_1$ . Therefore if any tuple  $t$  in the output of  $(E_1 \bowtie E_2)$  is filtered out by the selection of the left-hand side, all the tuples in  $E_1$  whose value is equal to  $t[E_1]$  are filtered out by the selection of the right-hand side. Therefore:

$$\forall t, t \notin \sigma_\theta(E_1 \bowtie E_2) \Rightarrow t \notin \sigma_\theta(E_1) \bowtie E_2$$

Using similar reasoning, we can also conclude that

$$\forall t, t \notin \sigma_\theta(E_1) \bowtie E_2 \Rightarrow t \notin \sigma_\theta(E_1 \bowtie E_2)$$

The above two equations imply the given equivalence.

This equivalence is helpful because evaluation of the right-hand side avoids producing many output tuples which are going to be removed from the result. Thus the right-hand side expression can be evaluated more efficiently than the left-hand side expression.

### 16.3

#### Answer:

- a.  $R = \{(1, 2)\}$ ,  $S = \{(1, 3)\}$   
The result of the left-hand side expression is  $\{(1)\}$ , whereas the result of the right-hand side expression is empty.
- b.  $R = \{(1, 2), (1, 5)\}$   
The left-hand side expression has an empty result, whereas the right hand side one has the result  $\{(1, 2)\}$ .

- c. Yes, on replacing the *max* by the *min*, the expressions will become equivalent. Any tuple that the selection in the rhs eliminates would not pass the selection on the lhs if it were the minimum value and would be eliminated anyway if it were not the minimum value.
- d.  $R = \{(1, 2)\}$ ,  $S = \{(2, 3)\}$ ,  $T = \{(1, 4)\}$ . The left-hand expression gives  $\{(1, 2, null, 4)\}$  whereas the right-hand expression gives  $\{(1, 2, 3, null)\}$ .
- e. Let  $R$  be of the schema  $(A, B)$  and  $S$  of  $(A, C)$ . Let  $R = \{(1, 2)\}$ ,  $S = \{(2, 3)\}$  and let  $\theta$  be the expression  $C = 1$ . The left side expression's result is empty, whereas the right side expression results in  $\{(1, 2, null)\}$ .

## 16.4

## Answer:

All the equivalence rules 1 through 7.b of section Section 16.2.1 hold for the multiset version of the relational algebra defined in Chapter 2.

There exist equivalence rules that hold for the ordinary relational algebra but do not hold for the multiset version. For example consider the rule :-

$$A \cap B = A \cup B - (A - B) - (B - A)$$

This is clearly valid in plain relational algebra. Consider a multiset instance in which a tuple  $t$  occurs 4 times in  $A$  and 3 times in  $B$ .  $t$  will occur 3 times in the output of the left-hand side expression, but 6 times in the output of the right-hand side expression. The reason for this rule to not hold in the multiset version is the asymmetry in the semantics of multiset union and intersection.

## 16.5

## Answer:

- The relation resulting from the join of  $r_1$ ,  $r_2$ , and  $r_3$  will be the same no matter which way we join them, due to the associative and commutative properties of joins. So we will consider the size based on the strategy of  $((r_1 \bowtie r_2) \bowtie r_3)$ . Joining  $r_1$  with  $r_2$  will yield a relation of at most 1000 tuples, since  $C$  is a key for  $r_2$ . Likewise, joining that result with  $r_3$  will yield a relation of at most 1000 tuples because  $E$  is a key for  $r_3$ . Therefore, the final relation will have at most 1000 tuples.
- An efficient strategy for computing this join would be to create an index on attribute  $C$  for relation  $r_2$  and on  $E$  for  $r_3$ . Then for each tuple in  $r_1$ , we do the following:
  - a. Use the index for  $r_2$  to look up at most one tuple which matches the  $C$  value of  $r_1$ .
  - b. Use the created index on  $E$  to look up in  $r_3$  at most one tuple which matches the unique value for  $E$  in  $r_2$ .

## 16.6

**Answer:**

The estimated size of the relation can be determined by calculating the average number of tuples which would be joined with each tuple of the second relation. In this case, for each tuple in  $r_1$ ,  $1500/V(C, r_2) = 15/11$  tuples (on the average) of  $r_2$  would join with it. The intermediate relation would have  $15000/11$  tuples. This relation is joined with  $r_3$  to yield a result of approximately 10,227 tuples ( $15000/11 \times 750/100 = 10227$ ). A good strategy should join  $r_1$  and  $r_2$  first, since the intermediate relation is about the same size as  $r_1$  or  $r_2$ . Then  $r_3$  is joined to this result.

## 16.7

**Answer:**

- a. Use the index to locate the first tuple whose *building* field has value “Watson”. From this tuple, follow the pointer chains till the end, retrieving all the tuples.
- b. For this query, the index serves no purpose. We can scan the file sequentially and select all tuples whose *building* field is anything other than “Watson”.
- c. This query is equivalent to the query:

$$\sigma_{\text{building} \geq \text{'Watson'} \wedge \text{budget} < 5000}(\text{department}).$$

Using the *building* index, we can retrieve all tuples with *building* value greater than or equal to “Watson” by following the pointer chains from the first “Watson” tuple. We also apply the additional criteria of *budget* < 5000 on every tuple.

## 16.8

**Answer:**

- a. First create relations  $r$  and  $s$ , and add some tuples to the two relations, before finding the plan chosen; or use existing relations in place of  $r$  and  $s$ . Compare the chosen plan with the plan chosen for a query directly equating  $r.A = s.B$ . Check the estimated statistics, too. Some databases may give the same plan, but with vastly different statistics.  
(On PostgreSQL, we found that the optimizer used the merge join plan described in the answer to the next part of this question.)
- b. To use hash join, hashing should be done after applying the `upper()` function to  $r.A$  and  $s.A$ . Similarly, for merge join, the relations should be sorted on the result of applying the `upper()` function on  $r.A$  and  $s.A$ . The hash or merge join algorithms can then be used unchanged.

16.9

**Answer:**

The above expressions are equivalent provided  $E_2$  contains only attributes  $A$  and  $B$ , with  $A$  as the primary key (so there are no duplicates). It is OK if  $E_2$  does not contain some  $A$  values that exist in the result of  $E_1$ , since such values will get filtered out in either expression. However, if there are duplicate values in  $E_2.A$ , the aggregate results in the two cases would be different.

If the aggregate function is min or max, duplicate  $A$  values do not have any effect. However, there should be no duplicates on  $(A, B)$ ; the first expression removes such duplicates, while the second does not.

16.10

**Answer:**

The interesting orders are all orders on subsets of attributes that can potentially participate in join conditions in further joins. Thus, let  $T$  be the set of all attributes of  $S1$  that also occur in any relation in  $S - S1$ . Then every ordering of every subset of  $T$  is an interesting order.

16.11

**Answer:**

FILL IN

16.12

**Answer:**

Each join order is a complete binary tree (every non-leaf node has exactly two children) with the relations as the leaves. The number of different complete binary trees with  $n$  leaf nodes is  $\frac{1}{n} \binom{2(n-1)}{(n-1)}$ . This is because there is a bijection between the number of complete binary trees with  $n$  leaves and number of binary trees with  $n - 1$  nodes. Any complete binary tree with  $n$  leaves has  $n - 1$  internal nodes. Removing all the leaf nodes, we get a binary tree with  $n - 1$  nodes. Conversely, given any binary tree with  $n - 1$  nodes, it can be converted to a complete binary tree by adding  $n$  leaves in a unique way. The number of binary trees with  $n - 1$  nodes is given by  $\frac{1}{n} \binom{2(n-1)}{(n-1)}$ , known as the Catalan number. Multiplying this by  $n!$  for the number of permutations of the  $n$  leaves, we get the desired result.

16.13

**Answer:**

Consider the dynamic programming algorithm given in Section 16.4. For each subset having  $k + 1$  relations, the optimal join order can be computed in time  $2^{k+1}$ . That is because for one particular pair of subsets  $A$  and  $B$ , we need constant time, and there are at most  $2^{k+1} - 2$  different subsets that  $A$  can denote. Thus, over all the  $\binom{n}{k+1}$  subsets of size  $k + 1$ , this cost is  $\binom{n}{k+1} 2^{k+1}$ . Summing

over all  $k$  from 1 to  $n - 1$  gives the binomial expansion of  $((1 + x)^n - x)$  with  $x = 2$ . Thus the total cost is less than  $3^n$ .

**16.14****Answer:**

The derivation of time taken is similar to the general case, except that instead of considering  $2^{k+1} - 2$  subsets of size less than or equal to  $k$  for  $A$ , we only need to consider  $k + 1$  subsets of size exactly equal to  $k$ . That is because the right-hand operand of the topmost join has to be a single relation. Therefore the total cost for finding the best join order for all subsets of size  $k + 1$  is  $\binom{n}{k+1}(k + 1)$ , which is equal to  $n\binom{n-1}{k}$ . Summing over all  $k$  from 1 to  $n - 1$  using the binomial expansion of  $(1 + x)^{n-1}$  with  $x = 1$  gives a total cost of less than  $n2^{n-1}$ .

**16.15****Answer:**

- a. The nested query is as follows:

```
select  S.account_number
from    account S
where   S.branch_name like 'B%' and
        S.balance =
        (select max(T.balance)
         from account T
         where T.branch_name = S.branch_name)
```

- b. The decorrelated query is as follows:

```
create table t1 as
    select branch_name, max(balance)
    from account
    group by branch_name
select  account_number
from    account, t1
where   account.branch_name like 'B%' and
        account.branch_name = t1.branch_name and
        account.balance = t1.balance
```

- c. FILL IN

- d. In general, consider the queries of the form:

```

select  ...
from     $L_1$ 
where    $P_1$  and
         $A_1$  op
        (select f( $A_2$ )
         from  $L_2$ 
         where  $P_2$ )

```

where  $f$  is some aggregate function on attributes  $A_2$  and  $op$  is some boolean binary operator. It can be rewritten as

FILL IN

